

Genetic Algorithms for Federated Learning: Two Distinct Approaches

GenFed vs. FedCSGA - Comprehensive Analysis

Sarthak Mishra | Abhimanyu Singh Rathore

CS6007 Project

November 1, 2025

Overview

- 1 Problem Statement
- 2 Literature Survey
- 3 Contribution / Proposed Method
 - Paper 1: GenFed (Zheng et al.)
 - Paper 2: FedCSGA (Wu et al.)
- 4 Theoretical Guarantees (if any)
- 5 Experimental Setup
- 6 Results & Analysis
 - Analysis of Source Papers
 - Our Re-Implementation Results
- 7 Conclusion & Future Work

Federated Learning: The Core Problem

Given N clients with local datasets \mathcal{D}_i , the goal is to solve:

$$\min_{w \in \mathbb{R}^d} f(w) = \sum_{i=1}^N \frac{n_i}{n} F_i(w) \quad (1)$$

where $F_i(w) = \frac{1}{n_i} \sum_{(x,y) \in \mathcal{D}_i} \ell(w; x, y)$ is the local loss.

Key Challenges Addressed by these Papers:

- **Data Heterogeneity (Non-IID):** Clients' data distributions \mathcal{D}_i vary significantly, causing local models to drift and slowing convergence.
- **System Heterogeneity:** Clients have different computation, memory, and network speeds.
- **Communication Bottlenecks:** Uploading models is slow; total training time is often constrained by a deadline.
- **Participant Reliability:** Some clients may be slow, drop out, or be malicious (Byzantine).

The Landscape of FL Client/Model Optimization

Our papers tackle the challenges of heterogeneity and efficiency. They fit into a broader landscape of FL optimization research:

Optimization-Based

- **Goal:** Maximize clients or minimize time.
- **FedCS:** Greedily selects clients who can meet a deadline.
- **Hybrid-FL:** Heuristic to select clients based on time and data distribution.
- **FedCSGA** Uses a full Genetic Algorithm to solve this NP-hard selection problem, outperforming simple greedy heuristics.

Utility-Function-Based

- **Goal:** Select clients with the highest “utility”.
- **Oort:** Defines utility by training loss and time, prioritizing clients with high impact and speed.
- **PyramidFL:** A fine-grained utility approach.
- **GenFed** Uses a GA-inspired “fitness” (validation accuracy) to select *models* post-training, not *clients* pre-training.

Two GA Approaches: When to Apply?

FedCSGA (Wu et al.) *Genetic Client Selection*

- **When:** Before training round.
- **What:** Selects the optimal set of clients to participate.
- **Goal:** Maximize number of clients that can finish within a time deadline T .
- **GA Use:** Full GA (crossover, mutation) to solve the complex, NP-hard scheduling problem.

GenFed (Zheng et al.) *Genetic Model Aggregation*

- **When:** After clients train.
- **What:** Selects the optimal set of models to aggregate.
- **Goal:** Maximize global model quality by filtering out low-performing or malicious models.
- **GA Use:** GA-inspired (fitness, selection) but only uses the “Selection” step.

GenFed: Genetic Mechanism Mapping

Idea: FL as Genetic Evolution

Genetic Mechanism	Federated Learning Component
Population	Set of client models $\{w_i\}$
Crossover	Server aggregation ($\sum w_i$)
Mutation	Local client training (Gradient Descent)
Fitness Evaluation	Model performance on a test set
Selection	Missing in traditional FL!

GenFed's Contribution: Add the missing **Selection** step.

- **How:** Server maintains a small, public, and privacy-safe **global validation set** \mathcal{D}_g^v .
- Before aggregation, the server evaluates all k received models:

$$\text{fitness}_i = \text{Accuracy}(w_i^t, \mathcal{D}_g^v)$$

- It **selects** only the **top- ρ_t** models with the highest fitness.
- It **aggregates** only these elite models to form w_{t+1} .

GenFed: Pseudocode (Algorithm 2)

```
1: Server-Side Execution:
2: Initialize  $w_g^0$ , Global validation set  $\mathcal{D}_g^v$ 
3: Initialization Phase: (Optional) Send supplemental data to clients with little data.
4: Training Phase:
5: for each round  $t = 1, 2, \dots, T$  do
6:    $S_t \leftarrow$  Select a subset of clients (e.g., 10 random clients)
7:   for each client  $k \in S_t$  in parallel do
8:     Send  $w_g^t$  to client  $k$ 
9:      $w_k^t \leftarrow \text{ClientUpdate}(w_g^t, \mathcal{D}_k)$  {Local training = Mutation}
10:    Send  $w_k^t$  back to server
11:   end for
12:   — GenFed Step —
13:   for each received model  $w_k^t$  do
14:      $\text{fitness}_k \leftarrow \text{Evaluate}(w_k^t, \mathcal{D}_g^v)$  {Evaluate fitness}
15:   end for
16:    $\rho_t \leftarrow \text{GetAggregationQuantity}(t, \rho_{\max}, \text{strategy})$ 
17:    $C_t \leftarrow$  Select top- $\rho_t$  models from  $S_t$  based on fitness {Selection}
18:    $w_g^{t+1} \leftarrow \sum_{i \in C_t} p_i w_i^t$  {Crossover (of elite models)}
19: end for
20: Return  $w_g^T$ 
```

FedCSGA: System Model & Problem

Goal: Maximize clients selected, given a hard deadline T .

Time Model:

- Computation: $t_i^{\text{comp}} = (e \cdot d_i) / c_i$ (epochs \times data / capacity)
- Communication: $t_i^{\text{comm}} = S / B_i$ (model size / bandwidth)
- Clients upload **sequentially** (one-by-one).

Key Insight: Computation can overlap with the *previous* client's communication. The total time depends on the **sequence \mathbf{q}** of clients.

Optimization Problem:

$$\max_{\mathbf{q}} |\mathbf{q}| \quad \text{s.t.} \quad \Theta_{|\mathbf{q}|}^{\mathbf{q}} \leq T$$

- $\mathbf{q} = \langle q_1, q_2, \dots \rangle$ is the client upload order.
- $\Theta_{|\mathbf{q}|}^{\mathbf{q}}$ is the total time for sequence \mathbf{q} (a complex function of overlaps).
- This is a variant of BIN PACKING, which is **NP-hard**.
- A simple greedy approach will fail; a GA is well-suited.

Chromosome: A sequence of clients, $\mathbf{q}_i = \langle q_{i1}, q_{i2}, \dots \rangle$.

Fitness Function (with Adaptive Penalty):

$$F(\mathbf{q}_i) = h(\mathbf{q}_i) - \lambda g(\mathbf{q}_i)$$

- **Objective** $h(\mathbf{q}_i)$: How “good” is the sequence?
 - **IID**: $h(\mathbf{q}_i) = |\mathbf{q}_i|$ (Just maximize client count)
 - **Non-IID**: $h(\mathbf{q}_i) = \sum_{j=1}^{|\mathbf{q}_i|} (1 - \alpha \cdot A(q_{ij}))$ (Balances count with accuracy A , prioritizing low-accuracy clients)
- **Penalty** $g(\mathbf{q}_i)$: How “infeasible” is it?
 - $g(\mathbf{q}_i) = e^{\max(0, (\Theta - T)/T)} - 1$ (Exponential penalty for exceeding deadline T)
- **Penalty Factor** λ :
 - $\lambda = \lambda_0 e^{\sqrt{r}}$ (where r = generation)
 - **Intuition**: Start with low penalty (explore), end with high penalty (enforce deadline).

FedCSGA: Pseudocode (Algorithm 1, Part 1/2)

```
1: Input: Client set  $M$ , deadline  $T$ , population size  $n$ , generations  $R$ 
2: Output: Optimal client sequence  $\mathbf{q}^*$ 
3: Initialization:
4:  $P_r \leftarrow \emptyset$  (population)
5: for  $i = 1$  to  $n$  do
6:    $k \leftarrow$  random client from  $M$ 
7:    $\mathbf{q}_i \leftarrow \langle k \rangle$ 
8:   Greedily add cheapest clients to  $\mathbf{q}_i$  until  $\Theta^{\mathbf{q}_i} > T$ 
9:   Add  $\mathbf{q}_i$  to  $P_r$ 
10: end for
```

FedCSGA: Pseudocode (Algorithm 1, Part 2/2)

```
12: Evolution:
13: while  $r < R$  do
14:    $P_{r+1} \leftarrow \emptyset$ 
15:   Crossover:
16:   for  $c = 1$  to  $n/2$  do
17:     Select  $\mathbf{q}_i, \mathbf{q}_j$  from  $P_r$  (e.g., tournament)
18:      $\mathbf{q}'_i, \mathbf{q}'_j \leftarrow \text{UniformCrossover}(\mathbf{q}_i, \mathbf{q}_j)$  with adaptive  $p_c$ 
19:   end for
20:   Mutation:
21:   for  $c = 1$  to  $n$  do
22:      $\mathbf{q}'_i \leftarrow \text{AdjacentSwapMutation}(\mathbf{q}_i)$  with adaptive  $p_m$ 
23:   end for
24:   Selection:
25:    $P_{r+1} \leftarrow \text{SelectBest}(P_r \cup P', n)$  using fitness  $F(\mathbf{q})$ 
26:    $r \leftarrow r + 1$ 
27: end while
28:  $\mathbf{q}^* \leftarrow$  Best feasible chromosome ( $\Theta \leq T$ ) from  $P_R$ 
29: Return  $\mathbf{q}^*$ 
```

GenFed (Zheng et al.):

- Contribution is conceptual (mapping FL to GA) , lacking a formal convergence proof.
- **Implied Guarantee:** Selecting models by validation accuracy biases the global model toward a better solution .
- This filtering also provides inherent robustness by discarding low-quality or Byzantine models

FedCSGA (Wu et al.):

- **NP-Hardness:** The paper formally states that the deadline-constrained, order-dependent client selection problem is **NP-hard**.
- This justifies using a meta-heuristic (like a GA) over a simple greedy algorithm (like FedCS), which is likely to get stuck in a local optimum.
- The GA is a tool to find a *high-quality approximate solution* to this NP-hard problem in polynomial time.

Base Training Configuration (Our Re-Implementation)

Model Architecture:

- MLP: 784 (input) \rightarrow 128 \rightarrow 64 \rightarrow 10 (output)
- Loss: CrossEntropyLoss
- Optimizer: SGD with momentum

Training Hyperparameters:

- Learning rate: $\eta = 0.01$
- Local epochs: $E = 5$
- Batch size: 32
- Device: CUDA/CPU auto-detect

Federated Setup:

- Datasets: MNIST
- Total Clients N : 10 (for our test) or 100 (in papers)
- Clients per Round k : 5 (our test) or 10 (in papers)
- Non-IID: Dirichlet distribution with α (lower α = more heterogeneous)

Key Hyperparameters Explained

GenFed (Zheng et al.):

- \mathcal{D}_g^v : **Global Validation Set**. A public, balanced dataset (e.g., 10% of MNIST test set) held by the server to evaluate model fitness.
- ρ_t : **Aggregation Quantity**. The number of “elite” models to select for aggregation.
- ρ_{\max} : **Max Aggregation**. The upper bound for ρ_t .
- **Strategy (Eq 5-9)**: The function (Linear, Sinusoidal, etc.) that governs how ρ_t changes over time t .

FedCSGA (Wu et al.):

- T : **Time Deadline**. The hard constraint (in seconds) that the entire round (computation + sequential upload) must fit into.
- n : **Population Size**. The number of chromosomes (sequences) in the GA (e.g., 90).
- R : **Generations**. The number of evolution steps the GA takes.
- p_c, p_m : **Crossover/Mutation Probabilities**. Dynamically adapted based on chromosome fitness to balance exploration/exploitation.
- α : **Non-IID Factor**. A value in $[0, 1]$ that balances maximizing client

Paper Results: GenFed (Zheng et al.)

Setup: 100 clients, 10 selected/round, Dirichlet $\alpha = 0.1$

Performance (vs. FedAvg)

- **MNIST:** 15 \times faster, +1.17%
Acc
- **SVHN:** 34 \times faster, +1.62%
Acc
- **Fashion:** 36 \times faster, +7.71%
Acc
- **CIFAR-10:** 16 \times faster, +3.98%
Acc

Byzantine Robustness (CIFAR-10)

Attack	FedAvg Drop	GenFed Drop
Label Flip	-5.90%	-3.01%
IPM	-2.93%	-0.77%
Mimic	-3.96%	-1.38%

Insight: The validation filtering naturally discards malicious models as “low fitness”, reducing attack impact by $\geq 50\%$.

Scalability

- FedAvg performance *degrades* as client count \uparrow
- GenFed performance *improves*

Paper Results: FedCSGA (Wu et al.)

Setup: 100 clients, GA ($n=90$, $R=10$), Non-IID ($\alpha = 0.7$)

IID Setting ($T = 3\text{min}$)

Method	Clients/Round	Accuracy
Fashion-MNIST		
FedAvg	2.3	88.7%
FedCS	7.8	89.1%
FedCSGA	8.5 (+9%)	89.9%
CIFAR-10		
FedAvg	1.6	70.6%
FedCS	6.5	72.3%
FedCSGA	7.0 (+8%)	76.0%

Non-IID Setting ($T = 5\text{min}$)


Method	Clients/Round	Accuracy
Fashion-MNIST		
FedAvg	3.0	57.8%
PyramidFL	9.0	66.1%
FedCSGA	13.5 (+50%)	66.5%
CIFAR-10		
FedAvg	2.9	38.8%
PyramidFL	7.5	53.3%
FedCSGA	11.2 (+49%)	54.0%

Key Insights:

- GA (FedCSGA) consistently finds better solutions (more clients) than the greedy (FedCS) or random (FedAvg) approaches.
- The Non-IID fitness function ($\alpha = 0.7$) is critical, dramatically improving accuracy over FedAvg (54.0% vs 38.8%) by balancing client count and data diversity.

Our Results: Accuracy & Loss

Setup: $N = 10$ clients, $k = 5$ selected, MNIST, $\alpha = 0.5$




Analysis:

- **GenFed (Green):** Achieved the highest accuracy, confirming the paper's claim that filtering models improves quality.
- **FedCSGA (Orange):** Underperformed.
- **FedAvg (Blue):** Baseline.

Method	Avg Accuracy
FedAvg (Baseline)	96.61%
GenFed ($\rho = 3$)	97.02% (+0.41%)
FedCSGA	96.18% (-0.43%)


Our Results: Client Selection Analysis



Analysis:

- **FedAvg (Left):** Selection is random, as expected. All 10 clients were selected roughly 20-30 times over 50 rounds.
- **FedCSGA (Center):** The GA was **highly** selective. It almost exclusively picked clients 0, 2, 4, 6, 8.
- **GenFed (Right):** Selection is random (like FedAvg) because GenFed selects **models** after training, it doesn't change the **client selection** policy.

Our Results: Timing Analysis (Extrapolated)



- **Time per Round (Left):** FedCSGA is stable; GenFed varies due to validation, FedAvg fluctuates moderately.
- **Time Ratios (Center):** GenFed slows with complexity; FedCSGA stays consistently fast.
- **Cumulative Time (Right):** FedCSGA is fastest overall; GenFed is slowest from validation overhead.

Lessons Learned: FedCSGA Failure Analysis

Root Cause: Stale Fitness Proxy

Problem in Our Implementation:

- FedCSGA's Non-IID fitness $h(\mathbf{q}) = \sum(1 - \alpha \cdot A(q_j))$ requires knowing client accuracy $A(q_j)$.
- We computed this accuracy **once** at round 0 (w_0).
- The GA used this *stale* accuracy as a proxy for all 50 rounds.
- As the global model w_t evolved, the round 0 accuracy became a meaningless proxy.
- The GA “overfit” to the stale proxy, picking clients (0,2,4,6,8) that were good at round 0, but not necessarily at round 40.

Proposed Fixes:

- 1 **Re-evaluate Fitness:** Re-run all clients on a validation set every ~ 5 rounds to get fresh accuracy data for the GA.
- 2 **Scale Up:** The GA's advantage over greedy is minimal at $N = 10$. The papers used $N = 100$, where the search space is vast and a GA is necessary.

GenFed vs. FedCSGA: Key Differences

Aspect	GenFed	FedCSGA
Stage	Post-training (Model Aggregation)	Pre-training (Client Selection)
Decision	Which <i>models</i> to aggregate	Which <i>clients</i> to select
Objective	Maximize model quality/fitness	Maximize clients under a time deadline
GA Use	Minimal (Fitness + Selection only)	Full (Crossover, Mutation, Selection)
Chromosome	N/A (simple ranking)	A <i>sequence</i> of clients $\langle q_1, \dots \rangle$
Fitness	Validation accuracy on \mathcal{D}_g^v	Function of $ \mathbf{q} $, time Θ , and accuracy $A(\mathbf{q})$
Key Tool	Global validation set	Time model + Adaptive GA operators
Solves...	Data heterogeneity, Byzantine attacks	System heterogeneity, Time constraints

Use **GenFed** when...

- ✓ You have a public, representative validation dataset.
- ✓ **Participant reliability** is a major concern (e.g., malicious clients, stragglers with buggy code).
- ✓ Your main goal is **fastest convergence** (fewer rounds) and highest final accuracy.
- ✗ You have no hard per-round time deadlines.

Use **FedCSGA** when...

- ✓ You have **strict time deadlines** per round (e.g., real-time applications).
- ✓ **System heterogeneity** (client speeds) is the main bottleneck.
- ✓ You can accurately profile client computation (c_i) and bandwidth (B_i).
- ✗ You do not have a public validation set to filter models.

Hybrid Approach: Use **FedCSGA** to select the max clients for a constraint T , then use **GenFed** to filter the received models.

Summary

- **GenFed:** A simple, powerful *model selection* filter.
 - 10-35 \times faster convergence (in rounds)
 - Inherent Byzantine robustness
 - Validated in our re-implementation
- **FedCSGA:** A powerful *client selection* scheduler.
 - Solves NP-hard time-constrained selection
 - Needs scale ($N \gg 10$) and fresh fitness data

Future Work

- **Our Project:**
 - Fix FedCSGA fitness staleness
 - Scale to $N = 100$ clients
 - Implement and test the Hybrid approach
- **Research:**
 - Multi-objective GA for FedCSGA (time, accuracy, *and* fairness)

<https://github.com/sihtsar/FedGA>

- ① Zheng, H., Chu, J., Li, Z., Ji, J., & Li, T. (2025). "Accelerating Federated Learning with genetic algorithm enhancements." *Expert Systems With Applications*, 281, 127636.
- ② Wu, J., Ji, H., Yi, J., & Liu, L. (2025). "Optimizing Client Selection in Federated Learning Base on Genetic Algorithm." *Cluster Computing*, 28, 400.
- ③ McMahan, B., et al. (2017). "Communication-Efficient Learning of Deep Networks from Decentralized Data." AISTATS.
- ④ Nishio, T., & Yonetani, R. (2019). "Client selection for federated learning with heterogeneous resources in mobile edge." ICC.
- ⑤ Lai, F., et al. (2021). "Oort: Efficient federated learning via guided participant selection." OSDI.

Thank You Questions?

Backup: GenFed ρ_t Scheduling

Goal: Find the best strategy for varying ρ_t (number of models to aggregate).

- **Strategy (1):** Constant $\rho_t = \rho_{\max}$. (Stable, but slow start)
- **Strategy (3):** Linear $\rho_t = \min(\rho_{\max} \cdot t/c + 1, \rho_{\max})$
- **Strategy (4):** Sinusoidal ($\pi/2$) $\rho_t = \min(\rho_{\max} \cdot \sin(t\pi/2c) + 1, \rho_{\max})$
- **Strategy (5):** Sinusoidal (π) (Unstable)

Result: Strategies (3) and (4) performed best.

- **Intuition:** Start with *small* ρ_t (e.g., $\rho_t = 1$) to aggregate only the **very** best model, allowing for rapid exploration.
- As training progresses, *increase* ρ_t to include more models, which provides stability and fine-tuning.

Backup: FedCSGA α Parameter Impact

Non-IID Fitness Function: $h(\mathbf{q}) = \sum_{j=1}^{|\mathbf{q}|} (1 - \alpha \cdot A(q_j))$

Table: Final Accuracy vs. α (Non-IID, $T = 5\text{min}$)

α	0.0	0.4	0.7	1.0
Fashion-MNIST	63.5%	65.2%	66.5%	66.1%
CIFAR-10	49.1%	52.2%	54.0%	53.3%

Analysis:

- $\alpha = 0$: Maximize clients only (same as IID). Ignores data heterogeneity.
- $\alpha = 1$: Only cares about accuracy. Selects too few clients, hurting diversification.
- $\alpha = 0.7$: **Optimal trade-off**. Balances selecting *many* clients (for diversity) with prioritizing clients that *need* training (low-accuracy clients).